

Kubernetes

page 1 of 2

1 Templating

Helm

Template-based package mgr for k8s. “Charts” are stored in a “repository.” Eg, ArtifactHub, a public analog to DockerHub. 3 private repos tools as well, such as Chart-Museum. Upon helm install, templates are resolved into k8s manifests & sent to the API server. CLI subcmds:

completion	create	dependency	env
get	history	install	lint
package	plugin	pull	repo
rollback	search	show	status
template	test	uninstall	upgrade
verify	version		

Kustomize

Preferred by k8s developers (and adopted into kubectl-proper) as a light-weight alternative to templates; kustomize is itself a k8s resource (per KRM) that layers k8s yaml Δs as code into subdirectories near the original manifests. View prospective Δs with kustomize build or apply them directly with kubectl apply -k.

2 Docker

Images are identified by their tag, which generally looks like: <ns>/<version>. For gcloud:

```
us-central1-docker.pkg.dev/  
<prjID>/<repoID>/<img>:<tag>
```

CLI

docker run <tag>	# local execution
docker build -t <name>	# from Dockerfile
docker images	# local imgs
docker ps	# running cntrs
docker inspect <cntr>	# detailed info
docker exec -it <cntr> bash	# bash in
docker stop <cntr>	# stop execut'n
docker rm <cntr>	# fully remove
docker tag <tag>	# identify
docker push <tag>	# upload to repo
docker login	# login to repo
docker commit	# cntr → cntr
docker logs	#

Dockerfile

```
FROM ubuntu:18.04 # initial image  
COPY . /app # move files in/out  
RUN make /app # configure  
CMD python /app/app.py # main command
```

Repositories

3 kubectl

Common Flags

--all-namespaces	#
--namespace <name>	#
--field-selector <spec>	#
--tail=<#>	#

-o <style>	# output, eg wide
-f <file>	# applies, etc
-c <cntr>	# specific cntr
--dry-run	# see test yaml
--pretty	# output format
--watch	# live Δs
--selector	# label selector

Basics

kubectl create <>	#
kubectl get <>	#
kubectl run <>	#
kubectl expose <>	#
kubectl delete <>	#
kubectl exec --stdin --tty demo.yaml	# shell
kubectl run <name> --image=	# image

Application Deployment

kubectl apply <>	#
kubectl annotate <>	#
kubectl autoscale <>	#
kubectl debug <>	#
kubectl diff <>	#
kubectl edit <>	#
kubectl kustomize <>	#
kubectl label <>	#
kubectl patch <>	#
kubectl replace <>	#
kubectl rollout <>	#
kubectl scale <>	#
kubectl set <>	#
kubectl wait <>	#

Debugging / Inspection

kubectl attach <>	#
kubectl auth <>	#
kubectl cp <>	#
kubectl describe <>	#
kubectl exec <>	#
kubectl logs <>	#
kubectl port-forward <>	#
kubectl proxy <>	#
kubectl top <>	#

Cluster Management

kubectl api-versions <>	#
kubectl certificate <>	#
kubectl cluster-info <>	#
kubectl cordon <>	#
kubectl drain <>	#
kubectl taint <>	#
kubectl uncordon <>	#

Kubectl Settings

kubectl alpha <>	#
kubectl api-resources <>	#
kubectl completion <>	#
kubectl config <>	#
kubectl explain <>	#
kubectl options <>	#
kubectl plugin <>	#
kubectl version <>	#

4 Namespaces

Namespace

spec (namespace spec)

finalizers	
status	(namespace status)
conditions	phase
status.conditions	(latest observations)
status	type lastTrans'nTime
message	reason

LimitRange

spec	(min/max resource usage)
limits	
spec.limits	(list of limit objs)
type	default def'tReq'st
max	maxLim.Req.Ratio min

ResourceQuota

spec	(aggregate hard limits per ns)
hard	scopeSelector scopes
status	(RQs enforce)
hard	used

5 Workloads

Pod

spec	(description of a pod)
containers	initContainers imagePullSec'ts
affinity	enableS'vcLinks nodeSelector
nodeName	run'tmClassN'me tolerations
schedulerName	topo.SpreadConst. prio.ClassName
volumes	act.DeadlineSecs restartPolicy
priority	setHostn.AsFQDN readinessGates
hostname	term.GracePer.Secs subdomain
dnsConfig	dnsPolicy hostNetwork
hostPid	hostIPC shareProc.Names.
svcAcc'tName	autom'dS.A.Token securityCont'xt
spec.container	(to run in a pod)
name	image im'gPullPolicy
command	args workingDir
ports	env envFrom
volumeMounts	vol.Devices resources
lifecycle	term.M'sgPath term.M'sgPolicy
livenessProbe	readinessProbe startupProbe
securityContext	stdin[Once] tty

spec.container.port	()
containerPort	hostIP hostPort
name	protocol
spec.container.env	(vars & their source)
name	value valueFrom
spec.container.volumeMounts	(into cntr)
mountpath	m'ntPropag't'n subPath
name	readOnly subPathExpr
spec.container.resources	(guidelines)
limits	requests
spec.container.lifecycle	(actions)
postStart	preStart
spec.container.securityContext	(config)
runAsUser	runAsNonRoot runAsGroup
readOnlyFS	procMount privileged
allowPriv.Escal'n	capabilities seccompProfile
seLinuxOptions	windowsOpt'ns

[Handler]	(actions to be taken)
exec	httpGet tcpSocket
[Handler.httpGet]	(request spec)
port	host httpHeaders
path	scheme
[[Node Pod PodAnti]Affinity]	(schedul'g)
preferredDuringSchedulingIgnoredDuringExecution	
requiredDuringSchedulingIgnoredDuringExecution	
probe	(aliveness health check)

exec	init'lDelaySec's	timeoutSec's
httpGet	term.GracePer.Sec's	periodSec's
tcpSocket	failureThresh.	successThresh.

status	(info about a pod)	
nom'ntedNodeNamePhase	startTime	
message	podIP[s]	reason
containerStatuses	hostIP	qosClass
initCont'rStatuses	conditions	
status.[init ephemeral]	containerStatuses	
name	image	imageID
containerID	state	lastState
ready	restartCount	started

ReplicaSet

Replaces deprecated “ReplicationController;” frequently generated by Deployments (rather than being created directly).

spec	(rc specification)
selector	template replicas
minReadySec's	
status	()
replicas	avail.Replicas fullyLabeledRep's
readyReplicas	conditions observedGenerat'n

Job

spec	(Job specification)
template	parallelism completions
compl'nMode	backoffLim't act'vDeadl'nSecs
ttlSecsAftFins'd	suspend selector
man'lSelector	
status	(job's current state)
startTime	compt'nTime active
failed	succeeded compl'tdIndexes
conditions	uncountedTerm'dPods
status.conditions	(latest observed state)
status	type lastProbeTime
message	reason lastTrans'nTime

CronJob

spec	(CJ specification)
jobTemplate	schedule startingDeadlineSecond
concurrencyPolicy	suspend success'lJobsHist'yLimit
status	(CJ current status)
active	lastScheduleTime lastSuccessfulTime

DaemonSet

spec	(DS specification)
selector	template minReadySec's
updateStrat.	revis'nHist'yLim't
status	(state of DS)
conditions	num'rAvail'ble curr.Num'rSched.
num'rReady	collisionNum'rCount updatedNum'rSched.
num'rUnavail.num'rMissched.	observedGenerat'n desiredNum'rSche.

StatefulSet

spec	(SS specification)
serviceName	replicas template
revis'nHist.Lim.	selector podM'gmtPolic.
vol.ClaimTempl'ts	updateStrat. minReadySec's
spec.updateStrategy	(for updating pods)
type	rollingUpdate
status	(current SS state)
currentReplicas	replicas readyReplicas
updatedReplicas	conditions collisionCount
obs'vdGeneration	avail.Replicas currentRevision
updatedRevision	

Deployment
spec (desired deployment behavior)
selector template replicas
minReadySeconds strategy revis'nHist.Limit
prog.Deadl'nSec's paused
spec.strategy (for replacing pods)
type rollingUpdate
status (most recently observed)
replicas avail.Replicas readyReplicas
unavail.Replicas updatedReplicas collisionCount
conditions obs'vdGeneration

6 Configuration

ConfigMap
[ConfigMap] (additional top-level keys)
binaryData data immutable

Secret
[Secret] (additional top-level keys)
stringData data type immutable

7 Network

Services expose pods behind 2 layers of indirect'n: 1) a Service url, which kube-dns resolves to a singular svc IP; 2) an EndPoints rsc that maps the svc IP to a list of pod IPs. Ingresses (L7) and LoadBalancers (L4) can expose svcs externally by redirecting requests made against a given public IP.

Often in the below, IPs & ports contain:
[IPAddress] (description of IP)
ip hostname nodeName targetRef
[Ports] (list of ports)
port protocol name appProtocol

Service
spec (attributes for creating)
selector loadBal.S'rcRanges type
ipFamilies pub.NotReadyAddr's clusterIP[s]
externalIPs h'lthCh'kNodeP'nt loadBal.IP
ports alloc.L.B.N'deP'rts exten.Name
ipFam'yPolicy internalTrafficPol. sessionAffin.
extern.Traff.Pol. sess.Affin.Config loadBal.Class
status (current service status)
conditions loadBalancer loadBal'r.ingress

Endpoint
IPs → svcs; often generated indirectly by svcs.
subsets (set of IPs, ports in a service)
addresses notReadyAddresses ports

Ingress
spec (desired ingress)
defaultBackend ingressClassName rules tls
spec.defaultBackend (how it is specified)
resource service
spec.rules (traffic routing)
host http
spec.tls (tls configuration)
hosts secretName
status.loadBalancer.ingress (status)
hostname ip ports

IngressClass
spec (ref → ingress controller)
controller parameters

spec.parameters (addit'nal ctrl'r config)
kind name apiGroup namespace
scope

NetworkPolicy
spec (allowed traffic from/to pods: default open)
podSelector policyTypes ingress egress
spec.ingress (inbound rules whitelisting)
from ports
spec.egress (outbound rules whitelisting)
to ports
spec.[in|e]gress.[from|to] (rule src/dest)
ipBlock n'nspscSelector podSelector

Nginx Ingress Controller
A default ingress controller, so, to install nginx:
kubectl apply -f URL/deploy.yaml
Extensive nginx controller config fields here.
To expose a service, create a normal ingress resource (as above) with an ingress class (shown here) to specify which ingress ctrl'r to use. (Or just assume the default.)
Customize the ingress' behavior with these nginx-specific annotations.

8 Security

ServiceAccount
[TopLevelFields] (additional top level fields)
autom'ntS.A.Token imagePullSecrets secrets

[Cluster]Role
rules (top-level set of PolicyRules)
apiGroups resources verbs
r'srcNames nonR'srcURLs

[Cluster]RoleBinding
roleRef (reference → role)
apiGroup kind name
subjects (ref → user ID)
apiGroup kind name namespace

CertificateSigningRequest
spec (cert request)
request signerName expir'tnSec's
extra groups uid
usages username
status (result of request)
certificate conditions
status.conditions (reasons for rejection/etc)
status.type lastTran'tnTime
message reason lastUpd'teTime

TokenRequest
spec (client parameters of request)
audiences boundObj.Ref expirationSec's
spec.boundObjRef (token applies to this)
apiVersion kind name uid
status ()
expirationTimestamp token

TokenReview
spec (attempt to authenticate)
audiences token
status (result of request)
audiences authentication error user

9 Storage

PersistentVolume
spec (spec of PV)
accessModes capacity claimRef
m'ntOptions nodeAffinity P.V.ReclaimPol.
stor.ClassName volumeMode hostPath
local

∃ many types of PVs, eg gcePersistentDisk.
status (current status of volume)
message phase reason

PersistentVolumeClaim
spec ()
accessModes selector resources
volumeName stor.ClassName volumeMode
dataSource dataSourceRef
status (current status, PVC)
accessModes capacity conditions phase

StorageClass
[TopLevel] (additional top-level fields)
provisioner allowVol.ExpansionAll'dTopologies
mountOptions parameters reclaimPolicy
vol.BindingMode

10 Cluster

HorizontalPodAutoscaler
spec (HPA spec)
maxReplicas minReplicas scaleTargetRef
targ.CPUUtil.%
status ()
currentReplicas desiredReplicas curr.CPUUtil.%
lastScaleTime obs'vdGeneration

PodDisruptionBudget
spec ()
maxUnavail. minAvail. selector
status (PDB status)
currentHealthy desiredHealthy disruptionsAll'd
expectedPods conditions disruptedPods
obs'vdGeneration

Node
spec ()
configSource externalID podCIDR[s]
providerID taints unschedulable
taints (affecting all pods on node)
effect key timeAdded value
status (current node status)
addresses allocatable capacity
conditions config daemonEndP'ts
images nodeInfo phase
vol'sAttached vol'sInUse

ComponentStatus
conditions (comp conditions observed)
status type error message

Event
[TopLevel] (additional top-level fields)
eventTime action deprecatedCount
dep.FirstT'stmp dep.Source reason
dep.LastT'stmp note regarding
report'gC'trllr relate rept'gInstance
series type

11 Kubernetes Internals

kubeadm
To build your own k8s cluster, install node machines / OSes, with relevant network config and container runtimes, then run:
kubeadm init # run on master node

Controllers (& k8s “Control Plane”)
These are a bevy of subscribers to API server, each lis-
tening for events of interest. They affect cluster Δs by
writing back to API server. Egs:
Replc'n Mgr ReplicaSet C. DaemonSet C.
StatefulSet C. Node C. Deployment C.
Job C. Service C. Endpoints C.
Namespace C. Persist.Vol. C.

Scheduler
Assigns pod → node by Δ-ing pod definition thru API
server, which then notifies kubelet of the Δ.

kubelet
Monitors API server for pods assigned to the local node;
then starts new pods' containers by invoking the cluster's
configured container runtime (most often Docker).

kube-dns
A pod that ∃ (∈ kube-system) to resolve cluster service
URLs to service IPs ∇ pods ∈ the cluster. DNS inquiries
are routed here by default; each each pod's /etc/res-
olv. conf file points to kube-dns.

kube-proxy
Each node has a kube-proxy that writes End-
Points contents to its node's iptables file.

API
API adheres to REST principals, employing HTTP (or
protobufs) to transfer data. REST's “resources” map to
k8s “objects” (eg, pods, namespaces, etc) via its “resource”
URL. kubectl CLI wraps HTTP calls. Most/all re-
sources have the following:
[top level] (top level fields)
apiVersion kind metadata spec status
[resource].metadata (generic, ∇ resource)
name generateName namespace
labels annotations finalizers
manag'dFields ownerRef's creat'nTim'p
del.Time'p del.Gr'cPer.Sec's generation
rs'ceVersion selfLink uid
[<rsrctype>List] (collect'n, <rsrctype>s)
apiVersion kind metadata items
Eg, for a ChronJob (other fields are obj-specific):
apiVersion: batch/v1; kind: ChronJob